

SARTHAK SHINDE

AI/ML Engineer — Pune, Maharashtra, India

shindesarthak36@gmail.com | +91 8767341151 | [linkedin.com/in/sarthak-shinde-799704259](https://www.linkedin.com/in/sarthak-shinde-799704259) | github.com/sarthak7-securtiy

Professional Summary

AI/ML Engineer with hands-on production deployment experience in Python, machine learning, and deep learning, not just coursework. Built and shipped an OCR ingestion pipeline and a real-time KPI reporting platform at Tata Motors, now used by 10+ stakeholders. Certified in Generative AI with LLMs and Applied Machine Learning (Coursera). Seeking a junior AI/ML or Data Science role to bring the same bias toward shipping working systems.

Technical Skills

Languages: Python, SQL

ML & AI: Supervised/Unsupervised Learning, Classification, Regression, K-Means Clustering, Ensemble Methods, Hyperparameter Tuning, Statistical Analysis, Model Deployment

GenAI & NLP: Generative AI, Prompt Engineering, LLM Fine-tuning, Hugging Face, Transformers, Sentiment Analysis, Text Classification

Computer Vision: OpenCV, Tesseract OCR, Image Preprocessing, Adaptive Thresholding, Morphological Transforms

Frameworks: TensorFlow, PyTorch, Scikit-learn, FastAPI, Flask, LangChain, SQLAlchemy

Data & Databases: Pandas, NumPy, PostgreSQL, ETL Pipelines, Feature Engineering, EDA

Visualization: Matplotlib, Seaborn, Plotly, Streamlit, Chart.js

DevOps & Cloud: Docker, CI/CD, GitHub Actions, Git, REST APIs, Google Colab, Jupyter/Jupyter Notebook, VS Code

Work Experience

AI/ML & Data Science Intern

Oct 2025 – Jan 2026

Tata Motors

Pune, Maharashtra, India

- Eliminated **100%** of manual data entry for **1,000+** scanned records, as measured by full automation of the intake process, by building an end-to-end OCR pipeline with **Tesseract** and **OpenCV** using adaptive thresholding and morphological preprocessing.
- Improved OCR text extraction accuracy by **25%**, as measured against baseline extraction rates, by designing a custom preprocessing chain (Gaussian blur, adaptive thresholding, morphological transforms) tuned across varied document types.
- Cut KPI reporting lag from weekly to real-time for **10+ stakeholders**, as measured by a **7-day** reduction in reporting turnaround, by deploying a Dockerized **FastAPI** platform exposing **REST APIs** in place of manual Excel workflows.
- Improved training effectiveness scores by **15%** across **5+** training cohorts, as measured by statistical trend analysis on cohort performance data, by performing exploratory data analysis and feature engineering using **Pandas** and **NumPy**.

Projects

No-Code ML Analytics Platform (AADES) | Python, Streamlit, Pandas, Scikit-learn, Plotly | Live Demo

- Built a no-code analytics platform automating full data science workflows—schema profiling, statistical visualizations, K-Means clustering, and linear regression—using a **Streamlit** frontend and **Scikit-learn**.
- Engineered interactive charts with **Plotly** (scatter matrices, residual plots, elbow curves) enabling non-technical users to extract insights from raw data uploads without writing code; deployed on **Streamlit Cloud**.

RAG Document Q&A System | Python, FastAPI, LangChain, Gemini, ChromaDB, Streamlit, Docker | Live Demo

- Built a full-stack RAG pipeline allowing users to upload multiple PDFs, generate embeddings via **Google Gemini** (**text-embedding-004**), and query document contents using natural language with exact page-level source citations.
- Designed a **LangChain** orchestration layer with **ChromaDB** vector store for semantic retrieval and conversational memory, cutting manual document lookup time by an estimated **80%** versus keyword-based search across large document sets.
- Deployed as a **Dockerized** multi-service application (FastAPI backend + Streamlit frontend) returning cited answers in under **3 seconds** per query, with a single-container entrypoint for portability and reproducible deployment.

Market Sentiment & Trading Performance Dashboard | Python, Scikit-learn, Streamlit, Plotly, NLP

- Built an NLP + K-Means clustering pipeline analyzing **10,000+** financial and sentiment records, segmenting traders into **3** behavioral archetypes (bullish, bearish, neutral) with **85%+** clustering silhouette-based consistency.
- Delivered a **Streamlit** + Plotly dashboard with a statistical correlation engine that translates sentiment patterns into actionable trading strategy recommendations across **10,000+** records in real time.

Education

Nutan College of Engineering & Research

Expected July 2026

B.Tech in Computer Science & Engineering

Pune, Maharashtra, India — CGPA: 7.00

Certifications & Courses

- Generative AI with Large Language Models — Coursera (2025)
- Applied Machine Learning — Coursera (2025)
- Large Language Models — Coursera (2025)
- Data Science Essentials with Python — Cisco Networking Academy (2026)
- Python Programming — HackerRank (2025)